

# **Interfaces for Efficient Software Composition on Modern Hardware**

Shoumik Palkar

Dissertation Defense

April 2, 2020



# Software composition: A mainstay for decades!

NOTES ON STRUCTURED PROGRAMMING

by

prof.dr.Edsger W.Dijkstra

Programming Techniques      R. Morris  
Editor

This paper discusses modularization as a mechanism for improving the flexibility and comprehensibility of a system while allowing the shortening of its development time. The effectiveness of a "modularization" is dependent upon the criteria used in dividing the system into modules. A system design problem is presented and both a conventional and unconventional decomposition are described. It is shown that the unconventional decompositions have distinct advantages for the goals outlined. The criteria used in arriving at the decompositions are discussed. The unconventional decomposition, if implemented with the conventional assumption that a module consists of one or more subroutines, will be less efficient in most cases. An alternative approach to implementation which does not have this effect is sketched.

Key Words and Phrases: software, modules, modularity, software engineering, KWIC index, software design

CR Categories: 4.0

Introduction

A lucid statement of the philosophy of modular programming can be found in a 1970 textbook on the design of system programs by Gouthier and Pont [1, ¶10.23], which we quote below:<sup>1</sup>

A well-defined segmentation of the project effort ensures system modularity. Each task forms a separate, distinct program module. At implementation time each module and its inputs and outputs are defined. The modules are interconnected via their interface with other system modules. At checkout time the integrity of the module is tested independently; there are few scheduling problems in synchronizing the completion of several tasks before checkout can begin. Finally, the system is maintained in modular fashion; system errors and deficiencies can be traced to specific system modules, thus limiting the scope of detailed error searching.

Usually nothing is said about the criteria to be used in dividing the system into modules. This paper will discuss that issue and, by means of examples, suggest some criteria which can be used in decomposing a system into modules.

PREPARATION OF PROBLEMS FOR EDVAC-TYPE MACHINES

JOHN W. MAUCHLY

ELECTRONIC CONTROL COMPANY

8.2.

On the Criteria To Be Used in Decomposing Systems into Modules


D.L. Parnas  
Carnegie-Mellon University

ment on the fact that there are now a ting machines. To Professor Aiken, as for large-scale machines for quite a w tures of this meeting. It shows how e earlier — "snowballing." Interest in ple learn of its potentialities.


in explanation of the phrase "EDVAC" stand for "Electronic Discrete Variable automatically from the computers now in ex Mark I, Mark II, the ENIAC, and the l EDVAC-type machines are now bei it is appropriate to discuss both the prep on of the machines for the problems. sider some of the fundamental characte nts which differ significantly from pres aring on the handling of problems: (i nctions, few in number, to which the tions as well as numerical quantities i ed in accordance with other instruc

Structure and Interpretation of Computer Programs

Harold Abelson and Gerald Jay Sussman with Julie Sussman





# The result? An ecosystem of libraries + users



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



NumPy



statsmodels



SciPy



Keras



P Y T O R C H



XGBoost

tensor





# Example: ML pipeline in Python



pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



NumPy



statsmodels



SciPy



Keras



P Y T Ø R C H



XGBoost

tensor



# Example: ML pipeline in Python

- + **Users can leverage** 1000s of expertly-developed libraries across many different domains



- On modern hardware, composition is **no longer a “zero-cost” abstraction**



# Example: the function call interface

Used to pass data between functionality via pointers to in-memory values.

```
void vdLog(float* a, float* out, size_t n) {  
    for (size_t i = 0; i + 8 < n; i += 8) {  
        __m256 v = _mm256_loadu_ps(a + i);  
        ...  
        _mm256_log2_ps(v, ...);  
        ...  
    }  
}
```



(1) Pass args through stack

Performance gap between these is growing!

# Example: composition with function calls


Growing gap between memory/processing speed makes function call interface worse!

```
// From Black Scholes  
// all inputs are vectors  
d1 = price * strike  
d1 = np.log2(d1) + strike
```



**Data movement** is often **dominant bottleneck** in composing existing functions

# Hardware Trends are Shifting Bottlenecks



1.  
2.

Kagi et al. 1996. Memory Bandwidth Limitations of Future Microprocessors. ISCA 1996  
McCalpin. 1995. Memory Bandwidth and Machine Balance in Current High Performance Computers. TCCA 1995.



# Do we need a new way to combine software?

- **Strawman: use a monolithic system**
  - "Legacy" applications: thousands of users of existing APIs
  - **Example:** Community of data scientists who use optimized Python libraries
- **Strawman: always use low-level languages (e.g., C++) or optimize manually**
  - Optimizations [still] require lots of manual work
  - **Example:** Manual optimizations in MKL-DNN



# Challenges for software composition today

**Research vision:** make software composition a zero-cost abstraction again!



# My Research: new interfaces to compose software on modern hardware

**Key idea:** Use *algebraic properties* of software APIs in *new interfaces* to enable new optimizations

## Examples of algebraic properties:


- $F()$ 's loops can be fused with  $G()$ 's loops
- $F()$ 's args can be split + pipelined with  $G()$
- $F()$  is parallelizable after externally splitting its args




# My Approach: Three interfaces with new systems to leverage their properties

Name	Interface/Properties	System
Weld	<b>Focus:</b> Data movement optimization and automatic parallelization over existing library APIs	
Split annotations		
Raw filtering	<b>Focus:</b> I/O optimization via data loading	

# Preview: What a new interface can achieve



Black Scholes model with Intel  
MKL: **3-5x** speedup with Weld  
and SAs



Querying 650GB of Censys JSON  
data in Spark: **4x** speedup with  
raw filtering



# Rest of this Talk

- Weld
- Split annotations
- Raw filtering
- Impact, open source, and concluding remarks





# Weld: A Common Runtime for Data Analytics

*CIDR '17*

*PVLDB '18*

**Shoumik Palkar**, James Thomas, Deepak Narayanan, Pratiksha Thaker, Rahul Palamuttam, Parimarjan Negi, Anil Shanbhag, Malte Schwarzkopf, Holger Pirk, Saman Amarasinghe, Samuel Madden, Matei Zaharia



# Motivation for Weld


- + Ecosystem of 100s of existing libraries and APIs
- Combining these libraries is no longer efficient!

**Example:** Normalizing images in NumPy + classifying them in with log. reg. in TensorFlow: **13x difference** compared to an end-to-end optimized implementation


Can we enable existing APIs to compose efficiently on modern hardware?



# Weld: A Common Runtime for Data Analytics




# Weld: A Common Runtime for Data Analytics



# Weld's Runtime API



# Runtime API uses lazy evaluation



# Weld's IR



# Weld IR: Expressing Computations

Designed to meet three goals:

## 1. Generality

support diverse workloads and nested calls

## 2. Ability to express optimizations

e.g., loop fusion, vectorization, and loop tiling

## 3. Explicit parallelism



# Weld IR: Internals

Small “functional” IR with two main constructs.

**Parallel loops:** iterate over a dataset

**Builders:** declarative objects to produce results

- *E.g.*, append items to a list, compute a sum
- Different implementations on different hardware
- **Read after writes: enables mutable state**

Captures relational algebra, functional APIs like Spark,  
linear algebra, and composition thereof



# Weld's Loops and Builders

## Example: Functional Operators

```
def map(data, f):  
    builder = new appender[T]  
    for x in data:  
        merge(builder, f(x))  
    result(builder)
```

Builder that appends items to a list.

```
def reduce(data, zero, func):  
    builder = new merger[zero, func]  
    for x in data:  
        merge(builder, x)  
    result(builder)
```

Builder that aggregates a value.




# Weld's Optimizer



# Optimizer Goal

Remove **redundancy** caused by composing independent libraries and functions.



# Removing Redundancy

Rule-based optimizations for removing redundancy in generated Weld code.

## Before:

```
tmp = map(data, |x| x * x)
res1 = reduce(tmp, 0, +)      // res1 = data.square().sum()
res2 = map(data, |x| sqrt(x))// res2 = np.sqrt(data)
```

## Each line generated by separate function.

- Unnecessary materialization of tmp
- Two traversals of data
- Vectorization? Output size inference?



# Removing Redundancy

Rule-based optimizations for removing redundancy in generated Weld code.

## Before:

```
tmp = map(data, |x| x * x)
res1 = reduce(tmp, 0, +)
res2 = map(data, |x| sqrt(x))
```

## After:

```
bld1 = new merger[0, +]
bld2 = new appender[i32]
          (len(data))
for x: simd[i32] in data:
    merge(bld1, x * x)
    merge(bld2, sqrt(x))
```



# Removing Redundancy

Rule-based optimizations for removing redundancy in generated Weld code.

## Before:

```
tmp = map(data, |x| x * x)
res1 = reduce(tmp, 0, +)
res2 = map(data, |x| sqrt(x))
```

## After:

```
bld1 = new merger[0, +]
bld2 = new appender[i32]
               (len(data))
for x: simd[i32] in data:
    merge(bld1, x * x)
    merge(bld2, sqrt(x))
```

Example: Loop Fusion Rule to Pipeline Loops



# Removing Redundancy

Rule-based optimizations for removing redundancy in generated Weld code.

## Before:

```
tmp = map(data, |x| x * x)
res1 = reduce(tmp, 0, +)
res2 = map(data, |x| sqrt(x))
```

## After:

```
bld1 = new merger[0, +]
bld2 = new appender[i32]
               (len(data))
for x: simd[i32] in data:
    merge(bld1, x * x)
    merge(bld2, sqrt(x))
```

Example: Vectorization to leverage SIMD in CPUs



# Results




# Partial Integrations with Several Libraries

**Libraries:** NumPy, Pandas, TensorFlow, Spark SQL



$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



Evaluated on 10 data science workloads  
+ microbenchmarks vs. specialized systems

# Weld Enables Cross-Library Optimization





Image whitening + logistic regression classification  
with NumPy + TensorFlow: **13x** speedup




# Weld can be integrated incrementally



**Benefits with incremental integration.**



# Weld enables high quality code generation



**SQL: Competitive with state-of-the-art and handwritten baseline (other benchmarks open source!)**



# Impact of Optimizations: 8 Threads

Experiment	All	-Fuse	-Unrl	-Pre	-Vec	-Pred	-Grp	-ADS	-CLO
DataClean	1.00	2.44	0.97	0.99	0.98	0.95			
CrimeIndex	1.00	195	2.04	1.00	1.02	0.96		3.23	
BlackSch	1.00	6.68		1.44	1.95		1.64		
Haversine	1.00	3.97		1.20	1.02				
Nbody	1.00	1.78		2.22	1.01				
BirthAn	1.00	1.02		0.97	0.98			1.00	
MovieLens	1.00	1.07		1.02	0.98			1.09	
LogReg	1.00	20.18		1.00				2.20	
NYCFilter	1.00	9.99		1.20	1.23	2.79			
FlightDel	1.00	1.27		1.01	0.96	0.96	5.50		1.47



All optimizations  
enabled.

More Impactful      Less Impactful



# Impact of Optimizations: 8 Threads

Experiment	All	-Fuse	-Unrl	-Pre	-Vec	-Pred	-Grp	-ADS	-CLO
DataClean	1.00	2.44	0.97	0.99	0.98	0.95			
CrimeIndex	1.00	195	2.04	1.00	1.02	0.96		3.23	
BlackSch	1.00	6.68		1.44	1.95		1.64		
Haversine	1.00	3.97		1.20	1.02				
Nbody	1.00	1.78		2.22	1.01				
BirthAn	1.00	1.02		0.97	0.98			1.00	
MovieLens	1.00	1.07		1.02	0.98			1.09	
LogReg	1.00	20.18		1.00				2.20	
NYCFilter	1.00	9.99		1.20	1.23	2.79			
FlightDel	1.00	1.27		1.01	0.96	0.96	5.50		1.47

More Impactful      Less Impactful



**Loop fusion:** Pipeline loops to reduce data movement.  
Up to **195x** difference



# Weld Prior Work

- Runtime code generation in databases
  - HyPer, LegoBase, DBLAB, Voodoo, Tupleware
  - Only target SQL or don't explicitly support parallelism
- Languages for parallel hardware
  - OpenCL, CUDA, SPIR, DryadLINQ, Spark, etc.
  - No effective cross-function optimization (even with LTO etc.)
- Monad comprehensions, Delite multiloops
  - Weld supports incremental integration, cross-library API, adaptive optimizations



# My Approach: Building three systems to leverage new interface properties

Name	Interface/Properties	System
Weld	IR to <b>extract parallel “structure”</b> of library functions	<b>Compiler</b> to enable data movement optimization + parallelization
Split annotations		
Raw filtering		





# **Split annotations: Optimizing Data-Intensive Computations in Existing Libraries**

*SOSP '19*

Shoumik Palkar and Matei  
Zaharia



# Problem with Compilers: Developer Effort

- Need to replace **every function** to use compiler intermediate representation (IR)
- IR **may not even support all optimizations** present in hand-optimized code

## Examples

Weld needs 100s of LoC to support NumPy, Pandas



 Closed    Numba compilation error #3293  
ajaychat3 opened this issue on Sep 7, 2018 · 2

TypingError  
<ipython-input-98-845f112395cc> in <m  
    30 param grid1=11

Tensorflow XLA makes it slower?  
Asked 2 years, 4 months ago   Active 2 years, 4 months ago   Viewed 569 times

I am writing a very simple tensorflow program with XLA enabled. Ba

1  

```
import tensorflow as tf  
  
def ChainSoftMax(x, n)  
    tensor = tf.nn.softmax(x)  
    for i in range(n-1):
```

**“Sorry, our compiler doesn’t recognize this pattern yet”**

**“Some ops are expected to be slower compared to hand-optimized kernels”**



# **Split Annotations (SAs)**

**Data movement optimizations and automatic parallelization on unmodified library functions**




# SAs Enable Pipelining + Parallelism

**Key idea:** split data to pipeline and parallelize it.



# SAs Enable Pipelining + Parallelism

**Without SAs:**




**d1 = price \* strike** ←  
d1 = np.log2(d1) + strike



# SAs Enable Pipelining + Parallelism

**Without SAs:**




$d1 = \text{price} * \text{strike}$

$d1 = \text{np.log2}(d1) + \text{strike}$  ←



# SAs Enable Pipelining + Parallelism

With SAs:




```
d1 = price * strike  
d1 = np.log2(d1) + strike
```



# SAs Enable Pipelining + Parallelism

With SAs:




```
d1 = price * strike  
d1 = np.log2(d1) + strike
```

Build execution graph, **keep data in cache** by passing cache-sized splits to functions.



# SAs Enable Pipelining + Parallelism

With SAs:



Collectively fit in cache

`d1 = price * strike`

`d1 = np.log2(d1) + strike`




Build execution graph, **keep data in cache** by passing cache-sized splits to functions.



# SAs Enable Pipelining + Parallelism

With SAs:



Collectively fit in cache

`d1 = price * strike`


`d1 = np.log2(d1) + strike` ←

Build execution graph, **keep data in cache** by passing cache-sized splits to functions.



# SAs Enable Pipelining + Parallelism

With SAs:




```
d1 = price * strike  
d1 = np.log2(d1) + strike
```

Build execution graph, **keep data in cache** by passing cache-sized splits to functions.



# SAs Enable Pipelining + Parallelism

With SAs:




```
d1 = price * strike  
d1 = np.log2(d1) + strike
```

Build execution graph, **keep data in cache** by passing cache-sized splits to functions.



# SAs Enable Pipelining + Parallelism

With SAs:



Parallelize over split pieces

Build execution graph, **keep data in cache** by passing cache-sized splits to functions.




# Example of a split annotation for MKL

```
@sa(n: SizeSplit(n, K), a: ArraySplit(n, K),  
     b: ArraySplit(n, K), out: ArraySplit(n, K))  
// Computes out[i] = a[i] + b[i] element-wise  
void vdAdd(int n, double *a, double *b, double *out)
```

## Benefits compared to JIT compilers:

- + No intrusive library code changes
- + Reuses optimized library function implementations
- + Does not require access to library code

# SAs can sometimes outperform compilers



Black Scholes using Intel MKL  
**5x speedups** by reducing  
data movement



# Challenges in designing SAs

1. Defining how to split data and enforcing **safe** pipelining
2. Building a lazy task graph **transparently**
3. Designing a **runtime** to execute tasks in parallel

# Challenges in designing SAs





1. Defining how to split data and enforcing **safe** pipelining
2. Building a lazy task graph **transparently**
3. Designing a **runtime** to execute tasks  in parallel

See paper for  
implementation details!



# How do SAs enforce safe pipelining?

E.g., preventing pipelining between matrix functions that iterate over row vs. over column:



**Okay to pipeline** – split matrix by row, pass rows to function.

**Cannot pipeline** – second function reads incorrect values.



# SAs use a type system to enforce safe pipelining

A **split type** uniquely defines how to split function arguments and return values.

```
@sa(n: SizeSplit(n, K), a: ArraySplit(n, K),  
     b: ArraySplit(n, K), out: ArraySplit(n, K))  
void vdAdd(int n, double *a, double *b, double *out)
```

# SAs use a type system to enforce safe pipelining

A **split type** uniquely defines how to split function arguments and return values.

```
@sa(n: SizeSplit(n, K), a: ArraySplit(n, K),  
     b: ArraySplit(n, K), out: ArraySplit(n, K))  
void vdAdd(int n, double *a, double *b, double *out)
```


**ArraySplit** depends on function arg. **n**, the **runtime size** of an array, and **K**, the **number of pieces**.



# Same split types = values can be pipelined

An SA defines a unique “splitting” for a value using a primitive called a **split type**.

```
@sa(n: SizeSplit(n, K), a: ArraySplit(n, K),  
     b: ArraySplit(n, K), out: ArraySplit(n, K))  
void vdAdd(int n, double *a, double *b, double *out)
```







**Same split types** enforce values split in the same way: we **can pipeline** if data between functions has matching split types.



# Example: Matrix Pipelining in NumPy

Split type for NumPy matrices encodes dimension + axis:

**MatrixSplit(Rows, Cols, Axis, K)**



**Split types match:** axis=0  
for both function calls


**Split types don't match:** axis=0  
for first call, axis=1 for second call



# How an annotator writes SAs

1. Define a split type (e.g., `ArraySplit`, `MatrixSplit`)
2. Write a **split function** and **merge function** for the type
3. Annotate functions using the defined split types

# Mozart: Our system implementing SAs




# Mozart: Our system implementing SAs

**In C++:** Memory protection for lazy evaluation


**In Python:** Meta-programming for lazy evaluation

**See paper for details!**



## Mozart Client Library

Builds a lazily evaluated task graph,  
determines when to execute it




## Mozart Runtime

Check + initialize split types, split  
data, execute functions in parallel



# Mozart: Our system implementing SAs



# Results


# Data Types and Libraries Demonstrated

**Libraries:** L1 + L2 BLAS (MKL), NumPy, Pandas, spaCy, ImageMagick



NumPy

pandas  
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$





**Data types and operators:** Arrays, Tensors, Matrices, DataFrame joins, grouping aggregations, image processing algorithms, functional operators (map, reduce, etc.)

# SAs require less integration effort than compilers


Library	#Funcs	LoC for SAs			Total	LoC for Weld		
		SAs	Split. API	Total		Weld IR	Glue	Total
NumPy	84	47	37	<b>84</b>	321	73	<b>394</b>	
Pandas	15	72	49	<b>121</b>	1663	413	<b>2076</b>	
spaCy	3	8	12	<b>20</b>				
MKL	81	74	90	<b>155</b>				
ImageMagick	15	49	63	<b>112</b>				



# SAs can match JIT compilers under existing APIs





nBody simulation: **4.6x**  
**speedup** over NumPy



pandas  
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$

Birth Analysis: **4.7x**  
**speedup** over pandas

# SAs can accelerate highly optimized libraries



Shallow Water eqn:  
**3x speedup** over MKL




Image filter: **1.8x speedup** over ImageMagick



# Across the 15 workloads we benchmarked:

SAs **perform within 1.2x of all compilers** in **nine** workloads

SAs **outperform all compilers** in **four** workloads

**Compilers outperform SAs by >1.2x** in **two** of our workloads

- Up to **6x slower**: This happens when code generation (e.g., compiling interpreted Python) matters

# SAs Prior Work

- Black box code generation interface + parallelization
  - Numba, Pydron, Dask, Ray, Cilk, OpenMP
  - No pipelining/cross-function optimizations, which is focus of SAs
- Vectorization and Batch Processing
  - X100, MonetDB, Spark SQL
  - SAs enable these for arbitrary black-box libraries rather than SQL
- Automatic loop tiling and loop optimizations
  - Scala Collections, Polyhedral model in LLVM, etc.
  - Found to be ineffective over black-box functions, no pipelining



# My Approach: Building three systems to leverage new interface properties

Name	Interface/Properties	System
Weld	IR to <b>extract parallel “structure”</b> of library functions	<b>Compiler</b> to enable data movement optimization + parallelization
Split annotations	Annotations to define <b>how to partition function inputs</b>	<b>Runtime</b> to pipeline data among <b>unmodified library functions</b>
		



# Raw filtering: Optimizing I/O pipelines by restructuring data loading

*PVLDB '18*

**Shoumik Palkar**, Firas  
Abuzaid, Peter Bailis, and  
Matei Zaharia



# Parsing: A Computational Bottleneck

Parse (  )




Today:  
parse full input → slow!



# Key Opportunity: High Selectivity

High selectivity especially  
true for **exploratory  
analytics**.



**40%** of customer Spark queries at Databricks **select < 20%** of data  
**99%** of queries in Censys **select < 0.001%** of data



**How can we exploit high selectivity to accelerate parsing?**



# Sparser: Filter Before You Parse


Sparser: **Filter before parsing** first using fast filtering functions with **false positives, but no false negatives**



Today:  
parse full input → slow!



# Results: Accelerating End-to-End Spark Jobs



Censys queries on 652GB of JSON data: up to **4x** speedup by using Sparser.



# My Approach: Building three systems to leverage new interface properties

Name	Interface/Properties	System
Weld	IR to <b>extract parallel “structure”</b> of library functions	<b>Compiler</b> to enable data movement optimization + parallelization
Split annotations	Annotations to define <b>how to partition function inputs</b>	<b>Runtime</b> to pipeline data among <b>unmodified library functions</b>
Raw filtering	Composable filters with false positives	<b>Library</b> for accelerating I/O of serialized data



# New composition interfaces can improve performance on modern hardware

- **Weld** used at NEC to support new vector accelerator, prototyped at Databricks, used in several labs



- Ongoing work at Stanford for extending **SAs** to bridge GPU and CPU libraries
- Teradata, Google have prototyped **raw filtering** internally



# Acknowledgements

# Acknowledgements

Thank you to my committee members!



Keith  
Winstein



Christos  
Kozyrakis



Mendel  
Rosenblum



John  
Duchi



# Acknowledgements

Thank you Matei for an inspiring graduate career!



# Acknowledgements

To FutureData, for great discussions, gossip, and friendships that I hope will last forever

Cody, Daniel, Deepti, Edward, Fiodar, Kaisheng, Keshav, Kexin, Peter Bailis, Peter Kraft, Pratiksha, Sahaana

To my office mates, for teaching me about sports, goofing off with me, and tolerating four years of terrible jokes

Deepak, Firas, James

To other friends who supported me outside of lab

Akshay, Aubhro, Jeff, Neil, Rohit, Stephanie, Sagar, Sahil, Yuval

And of course, to my wife Paroma, whose unwavering support made grad school one of the fondest times of my life, and the rest of my family: my parents Anjali and Prasad, my sister Ishani, my aunt and uncle Trupti and Sourja, and my two little cousins Shreya and Tvisha, all of who were collectively responsible for keeping me smiling for the last 26 years ☺



# Conclusion

**Thesis:** We can use *algebraic properties* of software APIs in *new interfaces* to enable new optimizations

Demonstrated with three interfaces/systems:

- **Weld**
- **Split Annotations**
- **Raw filtering**

shoumik@cs.stanford.edu

